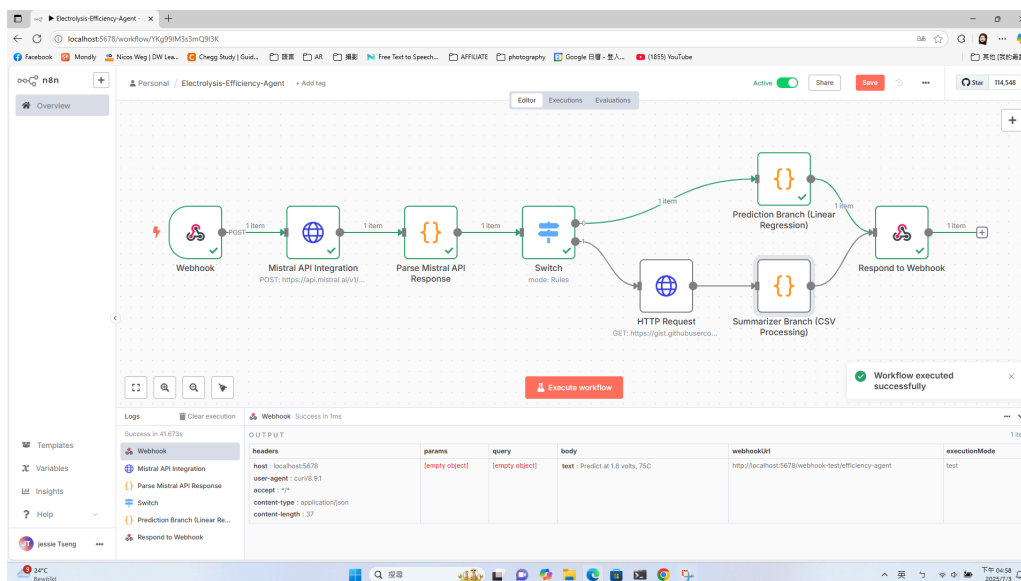


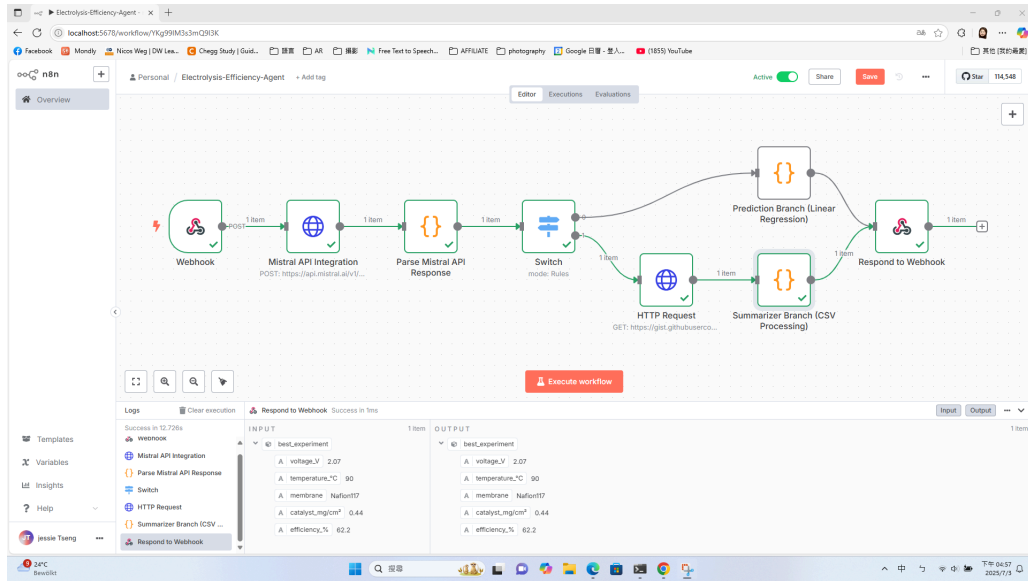
Electrolysis-Agent-Report

System Architecture

The system is built using **n8n** and consists of the following components:

1. **Webhook**– Accepts free-text requests from users.
2. **Mistral API Integration** – Sends the user's input to a hosted '**mistral-small-latest**' LLM model that interprets the request.
3. **Parse Mistral API Response** – Parses and validates the LLM's response into a clean JSON object
 - **intent** (either **"predict_run"** or **"summarize_best"**)
4. **Switch** – Directs the workflow based on the extracted **intent**.
5. **1.1 Prediction Branch** – Implements a pre-trained linear regression model with fixed coefficients.
 - 2.1 HTTP Requestre** – retrieves a public CSV file of 50 experimental runs.
 - 2.2 Summarizer Branch** – parses the CSV content by splitting the string into rows and columns, converting each row into a JSON object, identifying the row with the **highest efficiency_%**.
6. **Respond to Webhook** – Returns JSON output.





Linear Regression Model

1. Data Preprocessing Steps:

Step	Description
Data Source	CSV loaded from public GitHub Gist
Cleaning	Renamed columns, stripped whitespace, and dropped missing values
Features Used	voltage , temperature
Target Variable	efficiency
Scaling	Standardized features using StandardScaler
Train/Test Split	80% train, 20% test with random_state=42




Fig 1. Correlation Heatmap: choose the features


2. Model & Evaluation:

Step	Description
Model Type	Linear Regression inside a pipeline
Scaler Used	<code>StandardScaler</code> (removes mean and scales to unit variance)
Test RMSE	Computed using <code>mean_squared_error</code> on test set
Cross-Validation	5-fold <code>KFold</code> CV with <code>shuffle=True</code>
CV RMSE	Mean RMSE across all folds

$$\text{efficiency} = \beta_0 + \beta_1 \cdot \text{voltage} + \beta_2 \cdot \text{temperature}$$

Eq 1. LinearRegression equation, β_0 is the **intercept**, β_1 is the voltage **learned coefficient**, and β_2 is the temperature **learned coefficients**.

The model shown here uses voltage and temperature as inputs. Alternative models with more features (e.g. catalyst, membrane) were explored but not used in deployment.



- ✓ **Test RMSE:** ~1.04
- ✓ **5-Fold CV RMSE:** ~0.87

✓ **Final Regression Coefficients** for n8n:
`{'intercept': np.float64(22.1581), 'voltage': np.float64(10.2944), 'temperature': np.float64(0.1918)}`

Fig 2. Predicted vs Actual Plot

3. Exported Coefficients

These coefficients are recovered from the standardized model back to the **original scale**, suitable for integration:

```
const coefficients = {
  intercept: 22.1581,
  voltage: 10.2944,
  temperature: 0.1918
};
```

Example Inputs & Outputs

Prompt	Response
Predict run at 1.85V and 80C	{"prediction": "58.53%", "parameters": {"voltage": 1.85, "temperature": 80}}
Summarize the best efficiency run	{"best_experiment": {"voltage_V": 2.07, "efficiency_%": 62.2, ...}}
What is the expected efficiency at 1.7V and 60°C?	[{"prediction": "51.17%","parameters": {"voltage": 1.7,"temperature": 60},"model_info": "Linear Regression (voltage & temp)"}]