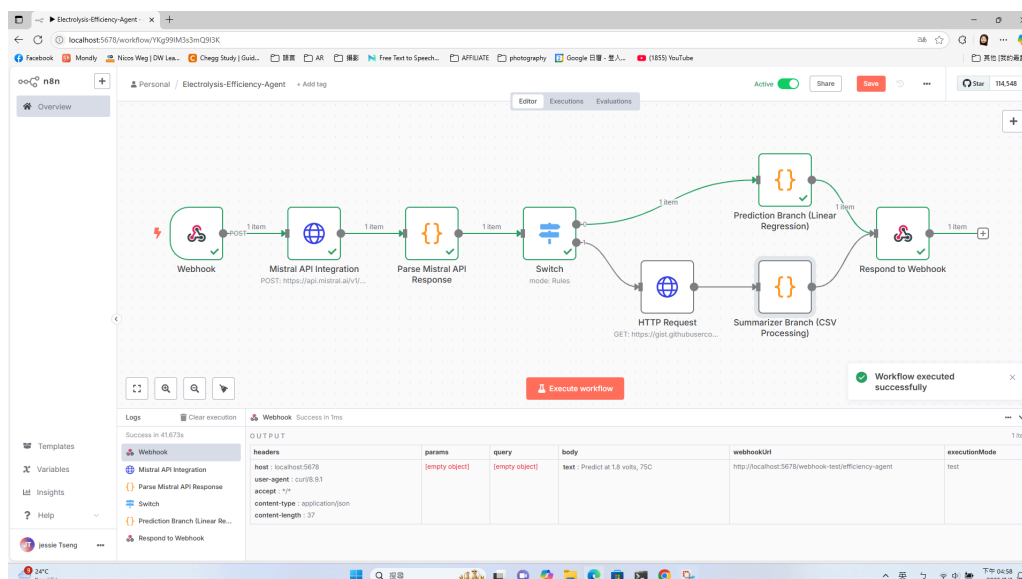


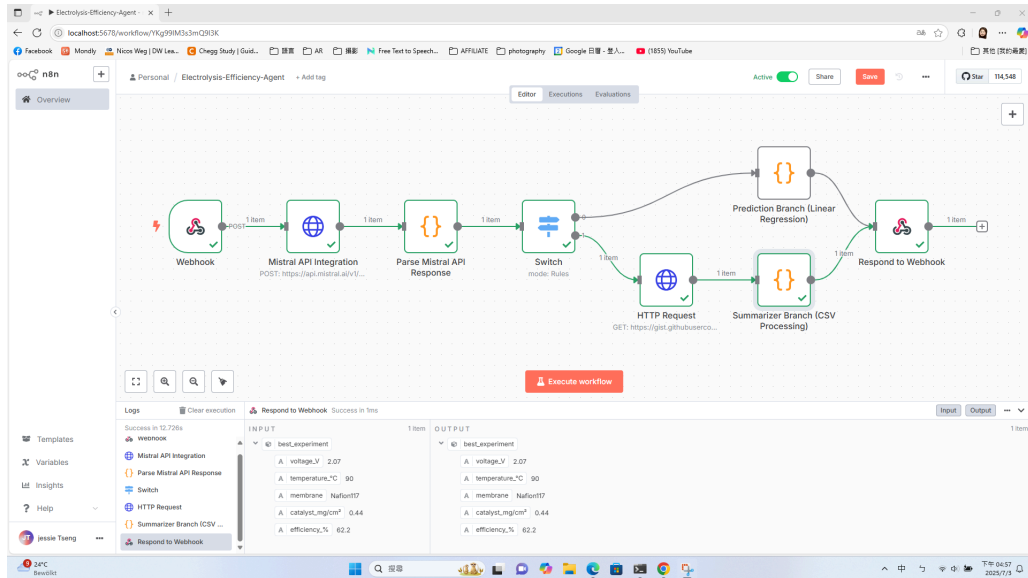
# Electrolysis-Agent-Report

## System Architecture

The system is built using **n8n** and consists of the following components:

1. **Webhook**– Accepts free-text requests from users.
2. **Mistral API Integration** – Sends the user's input to a hosted **LLM model** (e.g., Mistral) that interprets the request.
3. **Parse Mistral API Response** – Extracts structured information from the LLM's response, including:
  - **intent** (either **"predict\_run"** or **"summarize\_best"** )
  - relevant numeric parameters such as **voltage** and **temperature**
4. **Switch** – Directs the workflow based on the extracted **intent**.
5. **1.1 Prediction Branch** – Implements a pre-trained linear regression model with fixed coefficients.  
**2.1 HTTP Requestre** – retrieves a public CSV file of 50 experimental runs.  
**2.2 Summarizer Branch** – parses the CSV content by splitting the string into rows and columns, converting each row into a JSON object, identifying the row with the **highest efficiency\_%**.
6. **Respond to Webhook** – Returns JSON output to the user.





## Linear Regression Model

### 📌 Data Preprocessing Steps



1. **Source:** Dataset imported from external CSV URL.
2. **Column Standardization:** Renamed the following columns for consistency:
  - `voltage_V` → `voltage`
  - `temperature_°C` → `temperature`
  - `efficiency_%` → `efficiency`
  - `catalyst_mg/cm²` → `catalyst`
3. **Missing Data Handling:** Dropped all rows with missing values using `dropna()`.
4. **Feature Selection:**
  - Features used: `voltage`, `temperature`
  - Target variable: `efficiency`
5. **Train-Test Split:**
  - 80% training data, 20% test data
  - `random_state=42` for reproducibility
6. **Standardization:**

- Input features were standardized using `StandardScaler` to ensure zero mean and unit variance.

## 7. Regression Model:

- Applied `LinearRegression` from `scikit-learn` wrapped in a pipeline.

## Regression Model Performance

- **Model:** StandardScaler + LinearRegression
- **Test Set RMSE:**  
 `{rmse:.2f}`  
 (Root Mean Squared Error on hold-out test set)
- **Cross-Validation RMSE (5-Fold):**  
 `{cv_rmse.mean():.2f}`  
 (Mean RMSE across 5 shuffled splits using K-Fold CV)
- Model 1 (There is another model I try)

Test RMSE: 1.04  
5-Fold CV RMSE: 0.87

- **Coefficient Summary (Real Scale):**

Term	Coefficient
Intercept	22.1581
Voltage	10.294
Temperature	0.1918

## Example Inputs & Outputs

Prompt	Response
Predict run at 1.85V and 80C	<code>{"prediction": "58.53%", "parameters": {"voltage": 1.85, "temperature": 80}}</code>
Summarize the best efficiency run	<code>{"best_experiment": {"voltage_V": 2.07, "efficiency_%": 62.2, ...}}</code>
What is the expected efficiency at 1.7V and	<code>[{"prediction": "51.17%","parameters": {"voltage": 1.7,"temperature": 60},"model_info": "Linear Regression</code>

Prompt	Response
60°C?	(voltage & temp)"}]